

Titanic-Style Passenger Dataset | Kaggle Source

Minor Project Report

B.E. Artificial Intelligence & Data Science

East Point College of Engineering and Technology | VTU

Author: Zaid Khan | Academic Year: 2024-25

1. Project Overview

This report documents a complete End-to-End Exploratory Data Analysis (EDA) on a Titanic-style passenger dataset sourced from Kaggle. The objective was to demonstrate professional data science practices: auditing raw data quality, detecting and handling outliers using the IQR method, encoding categorical variables for Machine Learning readiness, and deriving meaningful insights through statistical visualization.

Dataset Information

Dataset Source	Kaggle — Titanic-Style Passenger Dataset
Kaggle Link	https://www.kaggle.com/competitions/titanic/data
File Format	.csv (Comma Separated Values)
Original Shape	305 rows x 8 columns (includes 5 duplicate rows)
Clean Shape	300 rows x 8 columns (after deduplication)
Key Features	PassengerId, Age, Fare, Gender, Embarked, Pclass, SibSp, Survived

Technical Workflow Summary

Step	Task	Tools / Methods
Step 1	Data Audit & Cleaning	<code>df.duplicated()</code> , <code>drop_duplicates()</code> , <code>fillna()</code>
Step 2	Outlier Detection (IQR)	<code>Q1</code> , <code>Q3</code> , <code>IQR</code> , <code>clip(lower, upper)</code>
Step 3	Encoding (Objects to Numerical)	<code>LabelEncoder</code> , <code>pd.get_dummies()</code>
Step 4	Visual Analysis	<code>Histogram</code> , <code>Scatter Plot</code> , <code>Box Plot</code> , <code>Heatmap</code>

2. Step 1 — Data Audit & Cleaning

The first step audits raw dataset quality by addressing two critical issues: **duplicate rows** and **missing values**. Duplicate entries cause model bias; missing values must be imputed to preserve dataset integrity.

2.1 Duplicate Detection & Removal

Cleaning Logic — Code:

```
dup_count = df.duplicated().sum() # Result: 5 duplicates found df = df.drop_duplicates() # Remove all exact duplicate rows print(df.shape) # Output: (300, 8)
```

Duplicates Found	5 exact duplicate rows
Action	Removed using df.drop_duplicates()
Shape Before	305 rows x 8 columns
Shape After	300 rows x 8 columns

2.2 Missing Value Handling

Column	Missing Count	Strategy	Reason
Age	20 nulls	Fill with Median (30.82)	Robust to outliers
Embarked	10 nulls	Fill with Mode ('S')	Most frequent category
Fare	0 nulls	No action needed	Column is complete
Gender	0 nulls	No action needed	Column is complete

Code:

```
df['Age'] = df['Age'].fillna(df['Age'].median()) # 30.82 df['Embarked'] = df['Embarked'].fillna(df['Embarked'].mode()[0]) # 'S'
```

3. Step 2 — Outlier Detection (IQR Method)

Outliers are extreme values that skew statistical measures and degrade model accuracy. The **Interquartile Range (IQR)** method was applied to **Age** and **Fare**. A **capping strategy** was chosen to preserve all 300 rows while bounding extremes.

IQR Formula

IQR = Q3 - Q1 (Q1 = 25th percentile, Q3 = 75th percentile)

Lower Bound = Q1 - 1.5 x IQR | Upper Bound = Q3 + 1.5 x IQR

Values outside [Lower, Upper] bounds are classified as outliers.

Action: **Cap** — `df[col].clip(lower, upper)` — limits values without data loss.

Outlier Logic — IQR Results Table

Column	Q1	Q3	IQR	Lower Bound	Upper Bound	Outliers
Age	22.24	37.65	15.41	-0.87	60.76	14 rows
Fare	9.74	54.84	45.10	-57.91	122.50	28 rows

Code:

```
for col in ['Age', 'Fare']: Q1, Q3 = df[col].quantile(0.25), df[col].quantile(0.75) IQR = Q3 - Q1 lower = Q1 - 1.5 * IQR upper = Q3 + 1.5 * IQR df[col] = df[col].clip(lower, upper) # CAP strategy
```




Figure 1: Box Plots for Age (left) and Fare (right) after IQR-based capping.

4. Step 3 — Encoding (Objects to Numerical)

Machine Learning models require all inputs to be numerical. Two object-type columns were identified and encoded using appropriate strategies based on their data nature.

4.1 Label Encoding — Gender Column

Label Encoding assigns integer labels to binary or ordered categories. Gender (male/female) is binary — ideal for this method. female=0, male=1.

```
from sklearn.preprocessing import LabelEncoder le = LabelEncoder() df['Gender_encoded'] = le.fit_transform(df['Gender']) # female -> 0 | male -> 1 # dtype changes from object -> int64
```

4.2 One-Hot Encoding — Embarked Column

One-Hot Encoding creates a binary column per category. Embarked has 3 unordered values (S=Southampton, C=Cherbourg, Q=Queenstown) — best handled with `pd.get_dummies()`.

```
df = pd.get_dummies(df, columns=['Embarked'], prefix='Emb') # Creates 3 new columns: Emb_C, Emb_Q, Emb_S # dtype changes from object -> bool
```

Encoding Proof — dtype Change

Before: Gender = *object (str)* | Embarked = *object (str)*

After: Gender_encoded = *int64* | Emb_C, Emb_Q, Emb_S = *bool*

All categorical columns successfully converted from object to numerical types.




Figure 2: Survival rate by Gender (left) and Embarked port distribution (right).

5. Step 4 — Visual Analysis

Four chart types were generated to analyze distribution, relationships, and correlations. All charts are clearly labeled with axes, titles, and legends.

5.1 Histogram — Age Distribution

After IQR capping and median imputation, age is approximately normally distributed. Mean = 30.86, Median = 30.82, Std = 12.33. The near-identical mean and median confirm minimal skew — a sign of successful outlier treatment.

5.2 Bar Chart — Survival Rate by Passenger Class

Survival rates across passenger classes were compared. The chart reveals the relationship between socioeconomic class and survival outcome — a key pattern in this dataset.




Figure 3: Histogram of Age distribution with mean/median lines (left) and Survival Rate by Pclass bar chart (right).

5.3 Scatter Plot — Age vs Fare (by Survival)

The scatter plot compares Age and Fare for survived (green) vs non-survived (red) passengers. Higher fare passengers appear more frequently in the survived cluster, visually reinforcing the fare-survival correlation.

5.4 Correlation Heatmap — All Numerical Features

The heatmap reveals pairwise correlations between all numerical columns. Fare shows positive correlation with Survival (0.030). Pclass shows negative correlation with Fare, meaning higher class = lower class number.




Figure 4: Scatter plot Age vs Fare colored by survival (left) and Correlation Heatmap (right).

6. Final Summary — 3 Significant Findings

Based on the complete EDA pipeline, three key statistically grounded findings were extracted that reveal patterns in passenger survival and data quality.

- 1

Fare Positively Correlates with Survival ($r = 0.030$)

Passengers who paid higher fares show a positive correlation with survival. This links to passenger class — higher fare indicates higher class, providing better access to lifeboats and evacuation positions on the ship.
- 2

Age Distribution Stabilized — No Significant Skew After Cleaning

After filling 20 missing Age values with median (30.82) and capping 14 outliers via IQR, the column stabilized: Mean = 30.86, Median = 30.82, Std = 12.33. The near-zero mean-median gap confirms successful outlier treatment.
- 3

Outlier Volume is High — 14% of Numeric Data Required Treatment

28 of 300 Fare rows (9.3%) exceeded the upper bound of 122.50, and 14 Age rows (4.7%) exceeded 60.76. Without IQR-based capping, these extremes would severely distort model training — making outlier handling a critical preprocessing step.

7. Evaluation Criteria — Checklist

Task	Requirement	Status
Duplicates	Are exact copies removed?	PASS — 5 rows removed via drop_duplicates()
Outliers	Is the IQR formula applied correctly?	PASS — Q1/Q3/IQR calculated, values capped
Encoding	Are all 'Object' columns converted?	PASS — Gender (Label), Embarked (One-Hot)
Visualization	Are charts clear & properly labeled?	PASS — Histogram, Scatter, Heatmap, Box Plot

8. Conclusion

This project successfully demonstrates all four pillars of a professional EDA pipeline. Starting from a raw Kaggle dataset with quality issues, the workflow systematically cleaned duplicates, handled missing values, detected and capped outliers using the statistically sound IQR method, converted categorical variables to machine-readable numerical formats, and produced clearly labeled visualizations revealing meaningful patterns. The cleaned and encoded dataset is now fully ready for Machine Learning tasks such as classification or regression modeling.

Author: Zaid Khan | **Dataset:** Titanic-Style Passenger Data | **Rows Cleaned:** 300 | **Outliers Capped:** 42 values | **Features Encoded:** 2 columns | **Charts:** 4 types | **Evaluation:** All 4 criteria PASS